

CS4287 Neural Computing

Lab 3: Week 5, S1 AY 24/25

Introducing TensorFlow

Preliminaries

Create an account in Google Colab

<https://colab.research.google.com>

Open a new file in Colab– a new Jupyter Notebook

- ➔ Jupyter Notebooks is an web-based interactive python development environment
- ➔ Colab provides a runtime with GPU support, useful for matrix computation.
- ➔ A matrix is a tensor!
- ➔ Tensorflow is a framework and APIs for deep learning

Basics of Jupyter Notebook:

- Type in a block consisting of 1 or more statements
- hit the “Run Cell”
- Hit “+Code” to enter a new block
- Hit “+Text” to enter comments

An example of a block could be L1-L3 in the Exercise 1 program

Exercise 1: A Neural Network in Keras and Tensorflow for MNIST

Based on: Antonio Gulli, Amita Kapoor, and Sujit Pal. *Deep Learning with TensorFlow 2 and Keras*, 2nd Edition. Packt> Press. 2019

Prog1 implements a neural network in TensorFlow 2.0 for the MNIST dataset

The code uses the Keras API which is a layer of abstraction on top of TensorFlow

The dataset is split into X_train for updating the weights, and X_test for accessing performance.

Data is converted into float32 to use 32 bit precision and normalised in the range [0..1]

The labels are loaded in Y_train and Y_test.

The network architecture is

Input = $28 * 28 = 784$ nodes

Hidden Layer: 10 nodes

Output layer: 1 node using the Softmax activation which outputs probabilities that sum to 1. It aggregates the 10 answers provided by the 10 neurons in the hidden layer.

Weights = $(784 * 10) + (10 * 1) = 7840 + 10 = 7850$

The cycle is epochs of training and validation. And when finished, testing.

```
L1.      import tensorflow as tf
L2.      import numpy as np
L3.      from tensorflow import keras

L4.      # for reproducibility
L5.      np.random.seed(1671)

L6.      # hyper-parameters
L7.      EPOCHS = 200
L8.      BATCH_SIZE = 128
L9.      VERBOSE = 1
L10.     NB_CLASSES = 10 # number of outputs = number of digits
L11.     N_HIDDEN = 128
L12.     VALIDATION_SPLIT=0.2 # how much TRAIN is reserved for VALIDATION

L13.     # loading MNIST dataset
L14.     # verify
L15.     # the split between train and test is 60,000, and 10,000 respectively
L16.     # one-hot coding is automatically applied
L17.     mnist = keras.datasets.mnist
L18.     (X_train, Y_train), (X_test, Y_test) = mnist.load_data()
L19.     print(X_train.shape[0], 'train samples')
L20.     print(X_test.shape[0], 'test samples')

L21.     #normalize in [0,1]
L22.     X_train, X_test = X_train / 255.0, X_test / 255.0
L23.     #X_train is 60000 rows of 28x28 values --> reshaped in 60000 x 784
L24.     RESHAPED = 784
L25.     #
```

```

L26.     X_train = X_train.reshape(60000, RESHAPED)
L27.     X_test = X_test.reshape(10000, RESHAPED)
L28.     Y_train = Y_train.astype('float32')
L29.     Y_test = Y_test.astype('float32')

L30.     #create the neural network
L31.     model = tf.keras.models.Sequential()
L32.     model.add(keras.layers.Dense(NB_CLASSES,
L33.         input_shape=(RESHAPED,)),
L34.         kernel_initializer='zeros',
L35.         name='dense_layer',
L36.         activation='softmax'))

L37.     # summary of the model
L38.     model.summary()

```

You must specify the optimiser used to update the weights. See https://www.tensorflow.org/api_docs/python/tf/keras/optimizers

Options include SGD, Adam, etc.

You must specify the Loss (error / cost / objective) function, see:

https://www.tensorflow.org/api_docs/python/tf/keras/losses

Common choices are:

- MSE: Mean Squared Error
- categorical_crossentropy: which defines the multiclass logarithmic loss. If the true class is c , and the prediction is y , then

$$L(c, p) = - \sum_i C_i \ln(p_i)$$

You must then specify the metrics used to evaluate performance, see

https://www.tensorflow.org/api_docs/python/tf/keras/metrics

Common choices are

- Accuracy: the % of correct predictions
- Precision
- Recall

```

L39.     # compiling the model
L40.     model.compile(optimizer='SGD',
L41.         loss='sparse_categorical_crossentropy',

```

```

L42.         metrics=['accuracy'])

L43.         #train the model
L44.         model.fit(X_train, Y_train,
L45.                 batch_size=BATCH_SIZE,
L46.                 epochs=EPOCHS,
L47.                 verbose=VERBOSE,
L48.                 validation_split=VALIDATION_SPLIT)

L49.         #evaluate the model
L50.         test_loss, test_acc = model.evaluate(X_test, Y_test)
L51.         print('Test accuracy:', test_acc)

```

Exercise 2: Add Hidden Layers

Add the following lines of code before the model declaration

```

L1. #One hot representation of the samples
L2. Y_train = tf.keras.utils.to_categorical(Y_train, NB_CLASSES)
L3. Y_test = tf.keras.utils.to_categorical(Y_test, NB_CLASSES)

```

Change the model as follows

```

L4. model = tf.keras.models.Sequential()
L5. model.add(keras.layers.Dense(N_HIDDEN,
L6.     input_shape=(RESHAPED,),
L7.     name='dense_layer', activation='relu'))
L8. model.add(keras.layers.Dense(N_HIDDEN,
L9.     name='dense_layer_2', activation='relu'))
L10. model.add(keras.layers.Dense(NB_CLASSES,
L11.     name='dense_layer_3', activation='softmax'))

```

Exercise 3

At the end of this lab please work through Colab basics:

https://colab.research.google.com/notebooks/intro.ipynb?utm_source=scs-index#scrollTo=GJBs_fIRovLc